RED QUEEN'S SYNC PROTOCOL FOR ETHEREUM

ANDREW ASHIKHMIN & ALEXEY AKHUNOV

Abstract. As of early 2019, it takes new Ethereum full nodes a few hours to synchronise the blockchain state using the eth/63 protocol Buterin et al. [2019]. We propose a new protocol and an algorithm for Ethereum snapshot sync that is faster and more robust in respect of the growing state size. The new protocol is also tailored towards the needs of light clients and allows data storage formats other than the canonical Merkle Patricia trie. Performance results from a model implementation are encouraging.

> "A slow sort of country!" said the Queen.
> "Now, here, you see, it takes all the running you can do, to keep in the same place. If you want to get somewhere else, you must run at least twice as fast as that!"
>
> Lewis Carroll, Through the Looking-Glass and What Alice Found There

1. Introduction

It currently takes new Ethereum full nodes a few hours to synchronise the blockchain state using the eth/63 protocol Buterin et al. [2019]. Moreover, the growing state size might potentially result in complete sync failure, as described in Akhunov [2019a].

As part of the Ethereum 1x effort, we propose a new sync protocol and algorithm, which we call Red Queen's. In our sync algorithm seeders reply with data as of their most recent block. That results in an inconsistent trie on the leecher initially ("phase 1"), which is patched later on ("phase 2"). The idea is similar to that of Leaf Sync (Swende [2019]). Other sources of inspiration include BitTorrent, Parity's Warp Sync Parity Technologies, and Firehose Sync Carver and Cloutier [2019].

Further, we strive to make the protocol work for light clients like Mustekala—see also Buterin et al. [2018].

N.B. In this document, we only discuss snapshot synchronisation rather than synchronisation from the Genesis block.

2. Notation

We mostly follow the conventions and notations of the Yellow Paper (Wood [2018]), for instance, $\mathbb{Y}$ denotes the set of nibble sequences. We use the letter π for prefixes of state or storage trie keys $\mathbf{k} \in \mathbb{B}_{32}$,

$$(1) \qquad \pi \in \mathbb{Y} \ \wedge \ ||\pi|| \leq 64$$

A key matches a prefix if and only if all their first nibbles are the same,

$$(2) \qquad \texttt{MATCH}(\mathbf{k}, \pi) \equiv \forall_{i<||\pi||} : \mathbf{k}'[i] = \pi[i]$$

($\mathbf{k}'$ is a sequence of nibbles, while $\mathbf{k}$ is a sequence of bytes.)

3. Protocol Specification

We propose the following new request/reply operative pairs[1]:

GetBytecode (0x20)
[reqID: $\mathbb{N}$, [codeHash$_0$: $\mathbb{B}_{32}$, codeHash$_1$: $\mathbb{B}_{32}$, ...]]

Request EVM code of smart contracts. Just like `GetNodeData` from the current version (eth/63) of Ethereum Wire Protocol Buterin et al. [2019], except:

(1) includes a request ID;
(2) will only return bytecode with the corresponding hash, not arbitrary node data.

Bytecode (0x21)
[reqID: $\mathbb{N}$, [code$_0$: $\mathbb{B}$, code$_1$: $\mathbb{B}$, ...]]

Reply to `GetBytecode`.

GetStateNodes (0x22)
[reqID: $\mathbb{N}$, stateRoot: $\mathbb{B}_{32}$, [prefix$_0$: $\mathbb{Y}$, prefix$_1$: $\mathbb{Y}$, ...]]

Request state trie nodes as of a specific state. Note that this operative is similar to `GetNodeData`, but it uses prefixes rather than hashes as node keys. It will also only return nodes from the state trie, not arbitrary node data. TODO: prefix–node correspondence is trivial for branch nodes, not so much for leaf or extension nodes. TODO: prefix encoding consistent with the Yellow Paper.

StateNodes (0x23)
[reqID: $\mathbb{N}$, [node$_0$: $\mathbb{B}$, node$_1$: $\mathbb{B}$, ...], [avalableState$_0$, ...]$_{opt}$]

Reply to `GetStateNodes`. The empty list $\varnothing$ returned instead of a node means that the peer does not have enough information about the node requested. TODO: available state roots.

GetAccounts (0x24)
[reqID: $\mathbb{N}$, stateRoot: $\mathbb{B}_{32}$, [prefix$_0$: $\mathbb{Y}$, prefix$_1$: $\mathbb{Y}$, ...]]

TODO: description.

Accounts (0x25)
[reqID: $\mathbb{N}$,
  [
    [status$_0$, [[key$_0^0$: $\mathbb{B}_{32}$, val$_0^0$: $\mathbb{B}$], [key$_0^1$: $\mathbb{B}_{32}$, val$_0^1$: $\mathbb{B}$], ...]$_{opt}$],
    [status$_1$, [[key$_1^0$: $\mathbb{B}_{32}$, val$_1^0$: $\mathbb{B}$], [key$_1^1$: $\mathbb{B}_{32}$, val$_1^1$: $\mathbb{B}$], ...]$_{opt}$],
    ...
  ],
  [avalableState$_0$, ...]$_{opt}$
]

Reply to `GetAccounts`.

GetStorageSizes (0x26)
[reqID: $\mathbb{N}$, stateRoot: $\mathbb{B}_{32}$, [addressHash$_0$: $\mathbb{B}_{32}$, addressHash$_1$: $\mathbb{B}_{32}$, ...]]

Request storage trie sizes as of a specific state.

StorageSizes (0x27)
[reqID: $\mathbb{N}$, [numLeaves$_0$: $\mathbb{N}|\varnothing$, numLeaves$_1$: $\mathbb{N}|\varnothing$, ...], [avalableState$_0$, ...]$_{opt}$]

---

[1]For some extra information see Péter Szilágyi's comment at ETH v64 Wire Protocol Ring.

Reply to `GetStorageSizes`. The peer may return the empty list $\varnothing$ instead of the number of leaves for accounts it does not have enough information about.

GetStorageNodes (0x28)
[reqID: $\mathbb{N}$, stateRoot: $\mathbb{B}_{32}$,
  [addressHash^0: $\mathbb{B}_{32}$, [prefix$_0^0$: $\mathbb{Y}$, prefix$_1^0$: $\mathbb{Y}$, ...]],
  [addressHash^1: $\mathbb{B}_{32}$, [prefix$_0^1$: $\mathbb{Y}$, prefix$_1^1$: $\mathbb{Y}$, ...]],
  ...
]

Request storage trie nodes as of a specific state.

StorageNodes (0x29)
[reqID: $\mathbb{N}$,
  [
    [node$_0^0$: $\mathbb{B}$, node$_1^0$: $\mathbb{B}$, ...],
    [node$_0^1$: $\mathbb{B}$, node$_1^1$: $\mathbb{B}$, ...],
    ...
  ],
  [avalableState$_0$, ...]$_{opt}$
]

Reply to `GetStorageNodes`. The empty list $\varnothing$ returned instead of a node means that the peer does not have enough information about the node requested.

GetStorageValues (0x2a)
[reqID: $\mathbb{N}$, stateRoot: $\mathbb{B}_{32}$,
  [addressHash^0: $\mathbb{B}_{32}$, [prefix$_0^0$: $\mathbb{Y}$, prefix$_1^0$: $\mathbb{Y}$, ...]],
  [addressHash^1: $\mathbb{B}_{32}$, [prefix$_0^1$: $\mathbb{Y}$, prefix$_1^1$: $\mathbb{Y}$, ...]],
  ...
]

Request storage subtrie leaves along with proof nodes as of block #`blockAtLeast` or newer. The empty list $\varnothing$ instead of the account hash signifies state rather than storage trie. `fromLevel` specifies the number of upper nodes to be excluded from the proof in case the chain has not moved ahead (reply block is not newer).

StorageValues (0x2b)
[reqID: $\mathbb{N}$,
  [
    [
      [status$_0^0$, [[key$_{00}^0$: $\mathbb{B}_{32}$, val$_{00}^0$: $\mathbb{B}$], [key$_{01}^0$: $\mathbb{B}_{32}$, val$_{01}^0$: $\mathbb{B}$], ...]$_{opt}$],
      [status$_1^0$, [[key$_{10}^0$: $\mathbb{B}_{32}$, val$_{10}^0$: $\mathbb{B}$], [key$_{11}^0$: $\mathbb{B}_{32}$, val$_{11}^0$: $\mathbb{B}$], ...]$_{opt}$],
      ...
    ],
    [
      [status$_0^1$, [[key$_{00}^1$: $\mathbb{B}_{32}$, val$_{00}^1$: $\mathbb{B}$], [key$_{01}^1$: $\mathbb{B}_{32}$, val$_{01}^1$: $\mathbb{B}$], ...]$_{opt}$],
      [status$_1^1$, [[key$_{10}^1$: $\mathbb{B}_{32}$, val$_{10}^1$: $\mathbb{B}$], [key$_{11}^1$: $\mathbb{B}_{32}$, val$_{11}^1$: $\mathbb{B}$], ...]$_{opt}$],
      ...
    ],
    ...
  ],
  [avalableState$_0$, ...]$_{opt}$
]

Reply to `GetStorageValues`. Returns subtrie leaves with proofs as of block #`blockNumber` $\geq$ `blockAtLeast`. If the peer does not have information regarding a particular subtrie, it should return the empty list $\varnothing$ (e.g. []) rather than [nodes, tooManyLeaves, leaves] for it. The nodes returned are the upper nodes of the trie down to the subtrie root, so that it is possible to verify that the leaves do belong to the Merkle Patricia trie in question. The first `fromLevel` upper nodes must be skipped if and only if `blockNumber = blockAtLeast`. (If `fromLevel = 0`, then the nodes must start with the root node.) `tooManyLeaves` is a boolean flag ($0 = false$, $1 = true$) indicating that the subtrie requested contains too many leaves. TODO: how many is too many? The leaves are represented as the list of their keys[2] and values. The peer may only return either all leaves of the subtrie or nothing. In the case of `tooManyLeaves` the leaves should not be returned[3]. Proof nodes must be sent in any case; they give us a means to detect faulty or malicious peers. Note that state trie replies do not inline storage tries, unlike Leaf Sync.

4. Suggested Sync Algorithm

Here we suggest a possible algorithm for full state and storage snapshot synchronisation using the protocol specified above; light clients are out of scope. We describe a modus operandi where the seeder replies with its most recent data, and the leecher has to handle trie data coming from different blocks. We suggest to perform synchronisation in two stages: during phase 1 the leecher obtains leaf data (with the necessary proof nodes) as of any reasonable block height, while during phase 2 it patches up the trie in order to catch up to the most recent block[4]. The idea was proposed in Swende [2019].

Let us focus on the state trie for the moment; we shall come back to storage sync later. For phase 1 we suggest sending `GetAccounts` requests with a single prefix per request, ditto for phase 2. All requested prefixes are of size d_1 during phase 1 and of d_2 during phase 2, $d_2 \geq d_1$. We elaborate on the values of d_1 and d_2 later. The leecher gradually builds the first upper d_2 levels of the Merkle Patricia trie[5]. (The full trie can be constructed if so desired, but only the upper d_2 levels are necessary for our algorithm.) Populated nodes are marked with the block number as of they are valid. The algorithm preserves the following invariant: parent's block is always no older than child's block.

During phase 1 the leecher requests each possible prefix of size d_1 exactly once (barring network failures and faulty peers). When sending a request, the leecher sets its `blockAtLeast` to the block of the root of the current (partially populated) trie, `fromLevel` to the number of populated nodes down the path/prefix that are of the same block as the root. Having received a reply, the leecher verifies its proof. If the proof is valid, the leecher writes received leaves to the database and updates the nodes along the prefix/path. By the end of phase 1, the leecher will have all accounts populated, albeit inconsistently.

Figure 1 shows an example of a phase 1 step with $d_1 = 3$ and $d_2 = 4$. Say the leecher is interested in prefix <7a3>. The trie on the left represents leecher's state before sending a request. Root's block is 5, so it sets `blockAtLeast = 5`. The leecher sets `fromLevel` to 1 since there is no need to re-send the root as part of the proof. It cannot set `fromLevel` higher as the other nodes along the path are older than the root and thus have to be refreshed. Suppose that the seeder replies with data as of a newer block #6. Since the block has changed, the seeder ignores `fromLevel` and sends full proof. The leecher saves received leaves to its database and updates the nodes (<>, <7>, <7a>, <7a3>). The result is displayed on the right of Figure 1.

At the beginning of phase 2, the leecher updates the trie in order to figure out which subtries have to be refreshed. For that, it uses the `GetNodeData2` operative. The leecher refreshes the trie level by level, starting from the root (level 0) and descending to the level $d_2 - 1$. Nodes at the same level may be requested in batch. Having refreshed nodes for a level, it knows which child nodes one level below

---

[2]It is feasible to return suffixes rather than full keys given that prefixes are known, but we deem the performance gain to be insignificant.

[3]In that case the peer must not return an empty list as that would imply that no leaves match the given prefix.

[4]The Red Queen's race is a nice metaphor for phase 2.

[5]d_2 is small enough so that we can reasonably assume that (almost) all nodes in question are branch nodes; see Akhunov [2019b].
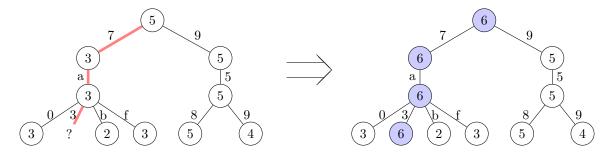
FIGURE 1. Illustration of phase 1 sync. Nodes are labelled with block numbers, and edges are labelled with nibbles.

have to be refreshed since the leecher can compare their hashes it currently holds with received fresh data. Therefore, only the nodes that have actually changed need to be requested. The leecher might have to restart the node refresh process from the root if new blocks are mined in between; however, provided a certain network bandwidth, the process converges. We analyse convergence conditions in the next section.

When all trie levels from the root down to the level $d_2 - 1$ are up to date, the leecher knows which nibbles at that last level have changed since phase 1. It refreshes the leaves corresponding to such nibbles using `GetAccounts` requests. That concludes the algorithm for state trie synchronisation.

The algorithm for storage sync is similar for large storage tries. Its parameters d_1 and d_2 are optimised based on the storage size as described in the next section. Small storage tries can be obtained in bulk requesting the empty prefix (meaning the entire trie) for a number of them in one go. `GetStorageSizes` provides a means of finding out storage sizes.

5. PERFORMANCE ANALYSIS

In this analysis, we assume that all tries are well balanced. We also assume that all top nodes up to a certain trie level i are branch nodes, not leaf nor extension nodes. This is a reasonable assumption if i is not too big—see Akhunov [2019b]. And we simplify the byte size function of the replies[6], ignoring overheads caused by auxiliary data such as `reqID`, RLP encoding, and the network layer. Let us introduce some notation:

n – the average node size in bytes, equal essentially to the size of a branch node as most nodes transferred will be branch nodes. 530 bytes is a good estimate.

l – the average leaf size in bytes, counting both key and value. For the state trie it is the average account size plus the size of its hash key, resulting in about 115 bytes. TODO: storage trie.

t – total number of leaves in a trie. For the state trie it is the number of accounts, which is about $53 \cdot 10^6$ as of February 2019—see Akhunov [2019b].

b – the network bandwidth available to the leecher.

τ – the block time, currently 15 seconds.

δ – the average number of leaf changes per block for a trie. For the state trie it is in the ballpark of 300.

$||R_n||$ – the number of nodes in a reply R.

$||R_l||$ – the number of leaves in a reply R.

We use the following simplified formula for the byte size of a reply R

$$(3) \qquad\qquad S(R) = ||R_n||n + ||R_l||l$$

The overhead of the sync algorithm during phase 1, compared with Parity's warp sync, is in the proof nodes sent alongside the leaf data. The overhead grows with d_1, so we want the trie depth to be as low as possible. On the other hand, small d_1 implies a large number of leaves per reply, which can be brittle or inefficient. Thus we set d_1 to the smallest value possible such that the replies are,

---

[6]Total request size is much smaller than total reply size, so we ignore requests as well.

on average, no larger than a certain size (say 32 KiB). We denote that maximum size as m. During phase 1 a `Accounts` reply contains at most d_1 nodes and its average number of leaves is $\frac{t}{16^{d_1}}$, which gives us

$$(4) \qquad d_1 n + \frac{t}{16^{d_1}} l \leq m$$

Fot the state trie the limit of 32 KiB yields $d_1 = 5$.

Let $C(d, \delta)$ be the maximum number of trie nodes from the upper d levels of a trie that can change (on average) per block[7]. At each level at most δ nodes can change, subject to δ being smaller than the number of nodes at the level. Thus

$$(5) \qquad C(d, \delta) = \sum_{i=0}^{d-1} \min(16^i, \delta)$$

If $16^2 \leq \delta \leq 16^3$ and $d \geq 3$, then

$$(6) \qquad C(d, \delta) = C'(d, \delta) \overset{\text{def}}{=} \delta(d - 3) + 273$$

We now analyse the minimum bandwidth required for the algorithm to converge during phase 2. At the very least, "to keep in the same place", we need to sync all changes per 1 block no slower than the block time τ. As previously described, the algorithm updates d_2 upper levels of the trie. So the upper bound on the number of nodes to be refreshed is $C(d_2, \delta)$. The number of subtries that need to be refreshed is no more than δ; each subtrie has $\frac{t}{16^{d_2}}$ leaves on average. Summing up, the total reply size per 1 block necessary not to lag behind is less than

$$(7) \qquad \texttt{RQS} \overset{\text{def}}{=} C(d_2, \delta) \, n + \delta \, \frac{t}{16^{d_2}} \, l$$

(RQS stands for Red Queen's Size). Though it is an upper bound, for our purposes `RQS` is close enough to the actual value. Differentiating, we find the value of d_2 that minimises `RQS`

$$(8) \qquad d_2^* = \frac{1}{\ln 16} \ln \left(\frac{tl \ln 16}{n} \right)$$

(Obviously, one has to round d_2^* up or down.) For the state trie the optimal $d_2^* = 6$ and the entailing `RQS` is about 0.7 MiB. Reiterating, the convergence condition for the state trie alone is

$$(9) \qquad b > \frac{\texttt{RQS}}{\tau}$$

For the Ethereum main net as of February 2019 this critical minimum bandwidth is about 0.4 Mbit/s. Table 1 shows performance results of an emulation of the sync protocol for various state trie sizes. The modelling code used is hosted at https://github.com/yperbasis/silkworm. Network latency was ignored in the emulation.

Bandwidth	10M	50M	100M
1 Mbit/s	03:39	18:44	39:04
10 Mbit/s	00:20	01:39	03:17
100 Mbit/s	00:02	00:10	00:20

TABLE 1. Emulated times of state trie sync for 10M, 50M, and 100M dust accounts.

Convergence analysis for large storage tries (e.g. CryptoKitties) is similar to the state trie analysis above.

6. CONCLUSION

TODO: conclusion.

---

[7]To be more precise mathematically, $C(d, \delta)$ is an upper bound on the expected value.

References

Alexey Akhunov. Looking back at the Ethereum 1x workshop 26–28.01.2019 (part 1), January 2019a. URL `https://medium.com/@akhounov/looking-back-at-the-ethereum-1x-workshop-26-28-01-2019-part-1-70c1ebd93266`.

Alexey Akhunov. Looking back at the Ethereum 1x workshop 26–28.01.2019 (part 2), February 2019b. URL `https://medium.com/@akhounov/looking-back-at-the-ethereum-1x-workshop-26-28-01-2019-part-2-d3d8fdcede10`.

Vitalik Buterin et al. Light Client Protocol, 2018. URL `https://github.com/ethereum/wiki/wiki/Light-client-protocol`.

Vitalik Buterin, Gavin Wood, et al. Ethereum Wire Protocol (ETH), 2019. URL `https://github.com/ethereum/devp2p/blob/master/caps/eth.md`.

Jason Carver and Brian Cloutier. Firehose Sync v0.2, April 2019. URL `https://notes.ethereum.org/Ov_W4E8lROazqYymPAF7Ew`.

Parity Technologies. Warp Sync Snapshot Format. URL `https://wiki.parity.io/Warp-Sync-Snapshot-Format`.

Martin Holst Swende. Leaf sync, March 2019. URL `https://notes.ethereum.org/kphcc_CKT4a5sUs_zWVelA`.

Gavin Wood. Ethereum: A Secure Decentralised Generalised Transaction Ledger, December 2018. URL `https://github.com/ethereum/yellowpaper`.